

# Lecture 10: Ultralytics Vision AI

[https://github.com/kaopanboonyuen/SC310005\\_ArtificialIntelligence\\_2025s1](https://github.com/kaopanboonyuen/SC310005_ArtificialIntelligence_2025s1)

---

Teerapong Panboonyuen  
<https://kaopanboonyuen.github.io>

# References

- <https://www.ultralytics.com/>
- <https://labelstud.io/>
- <https://roboflow.com/>
- <https://cs231n.stanford.edu/index.html>
- <https://encord.com/blog/object-detection/>
- <https://neptune.ai/blog/object-detection-algorithms-and-libraries>
- <https://learnopencv.com/non-maximum-suppression-theory-and-implementation-in-pytorch/>





Register now for  
 ultralytics  
YOLO Vision



ultralytics  
YOLO Vision



Products Solutions Resources Licensing Company

 GitHub ★ 113,533

Get in touch

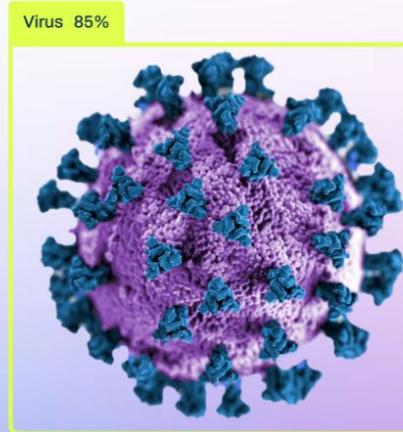
 EN

# Turn images into AI to get useful insights with no code

Drag-and-drop your data, do some tweaks — and voilà, a new powerful AI tool is born

Start for free

GitHub



Ask AI 

X

Register now for



→

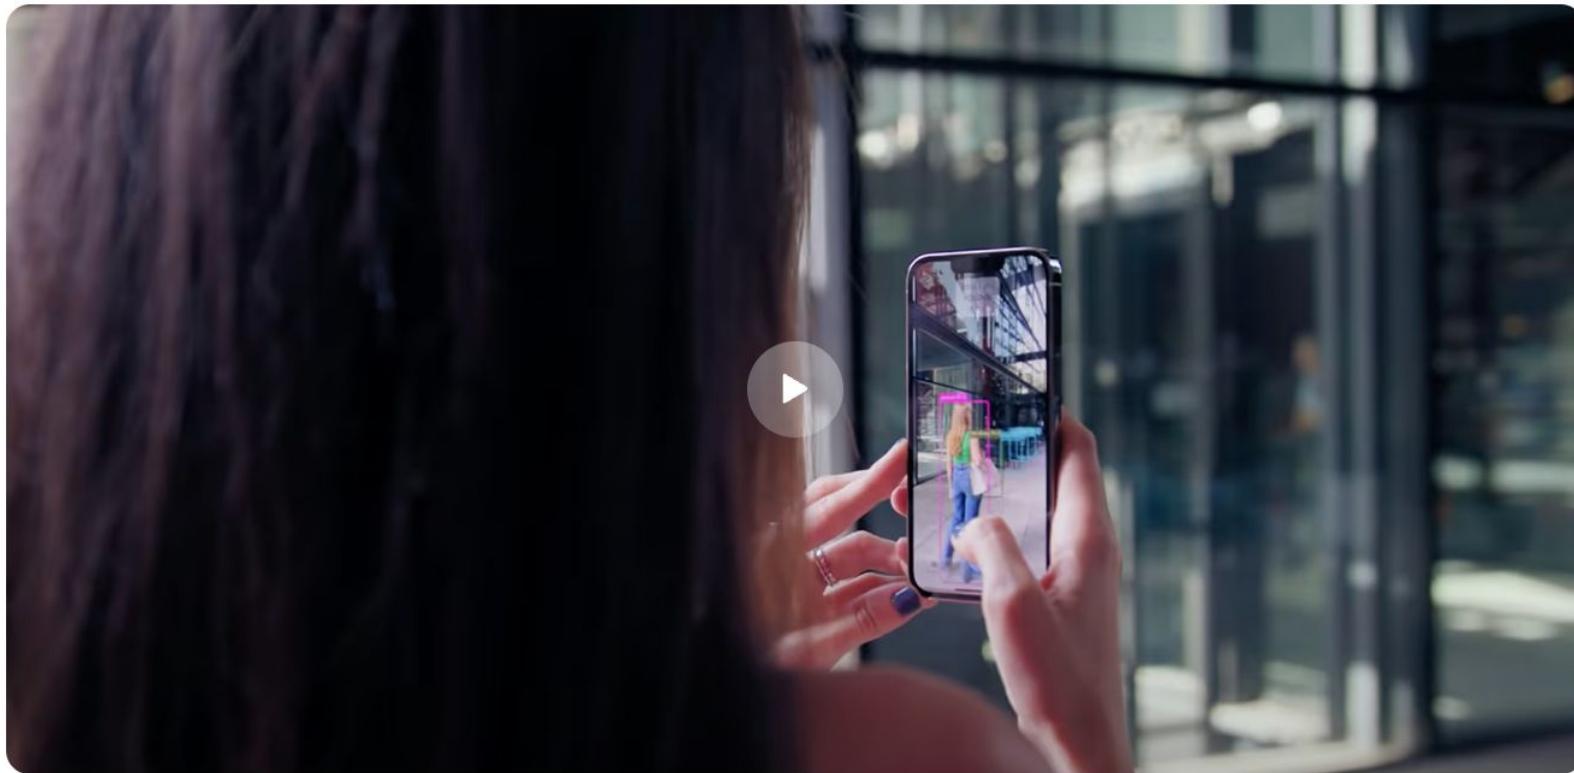


Products   Solutions   Resources   Licensing   Company

GitHub ★ 113,533

Get in touch

EN



[X](#)

Register now for  ultralytics  
YOLO Vision



ultralytics  
YOLO Vision



Products Solutions Resources Licensing Company

 GitHub ★ 113,533

Get in touch

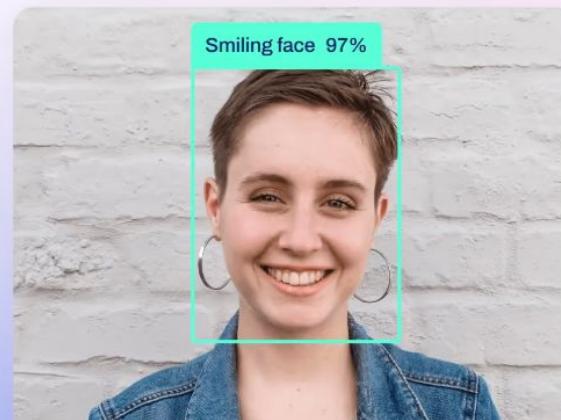
 EN

# Train computer vision models in seconds with Ultralytics YOLO

Explore our state-of-the-art AI architecture to train and deploy your highly-  
accurate AI YOLO models like a pro

Get started

[GitHub](#)



Faces

7 classes

7 785 img

6.7 GB



**Models**

YOLOv3

YOLOv4

YOLOv5

YOLOv6

YOLOv7

YOLOv8

YOLOv9

YOLOv10

YOLO11  NEW

YOLO12

SAM (Segment Anything Model)

SAM 2 (Segment Anything Model 2)

MobileSAM (Mobile Segment Anything Model)

FastSAM (Fast Segment Anything Model)

YOLO-NAS (Neural Architecture Search)

RT-DETR (Realtime Detection Transformer)

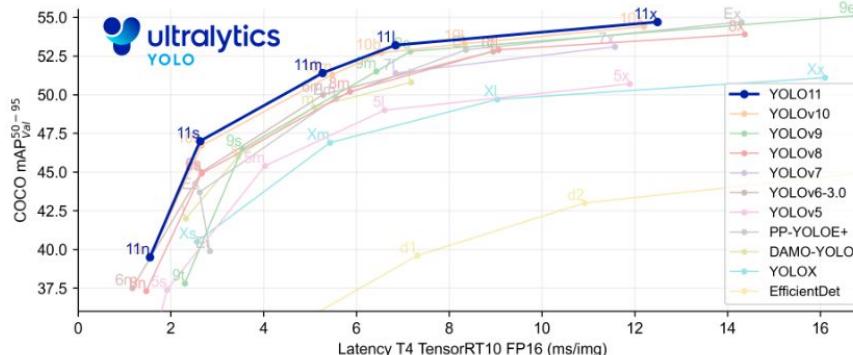
YOLO-World (Real-Time Open-Vocabulary Object Detection)

# Ultralytics YOLO11



## Overview

YOLO11 is the latest iteration in the Ultralytics YOLO series of real-time object detectors, redefining what's possible with cutting-edge accuracy, speed, and efficiency. Building upon the impressive advancements of previous YOLO versions, YOLO11 introduces significant improvements in architecture and training methods, making it a versatile choice for a wide range of computer vision tasks.

**Table of contents**[Overview](#)[Key Features](#)[Supported Tasks and Modes](#)[Performance Metrics](#)[Usage Examples](#)[Citations and Acknowledgements](#)[FAQ](#)

What are the key improvements in Ultralytics YOLO11 compared to previous versions?

How do I train a YOLO11 model for object detection?

What tasks can YOLO11 models perform?

How does YOLO11 achieve greater accuracy with fewer parameters?

Can YOLO11 be deployed on edge devices?

**Ask AI** 

---

## **Introduction to Computer Vision**

- **Definition:** Computer Vision enables machines to “see” and understand the visual world from images and videos.
  - **Goal:** Extract meaningful information for analysis, decision-making, and automation.
  - **Applications:** Healthcare, autonomous vehicles, wildlife monitoring, manufacturing, security, AR/VR.
-

## Types of Computer Vision Tasks

- **Image Classification:** Assign a label to an entire image (e.g., cat vs dog).
- **Object Detection:** Locate and classify objects within an image with bounding boxes.
- **Segmentation:**
  - **Semantic Segmentation:** Classify each pixel into a class.
  - **Instance Segmentation:** Detect objects and separate each instance.
- **Pose Estimation:** Detect keypoints on humans, animals, or objects to understand posture or movement.
- **Modern/Advanced AI Vision:** Generative AI (e.g., DALL-E, MidJourney, Text-to-Image) – will cover in advanced lectures.

# Other Computer Vision Tasks

Semantic Segmentation



GRASS, CAT,  
TREE, SKY

No objects, just pixels

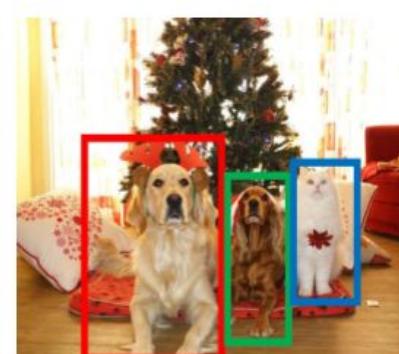
Classification + Localization



CAT

Single Object

Object Detection



DOG, DOG, CAT

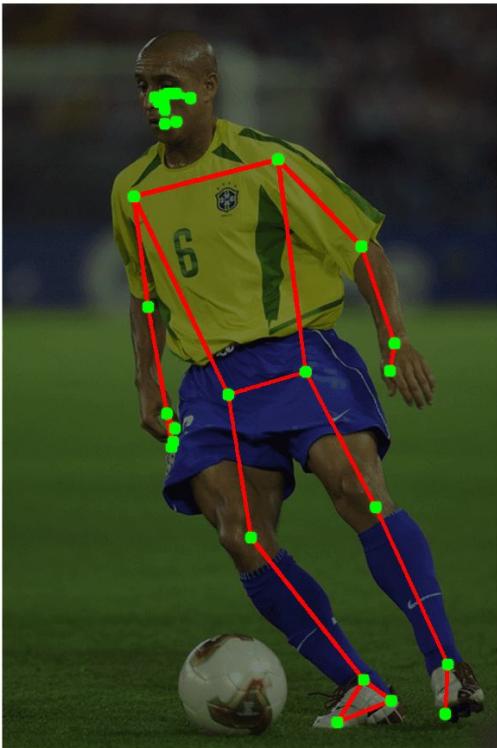
Multiple Object

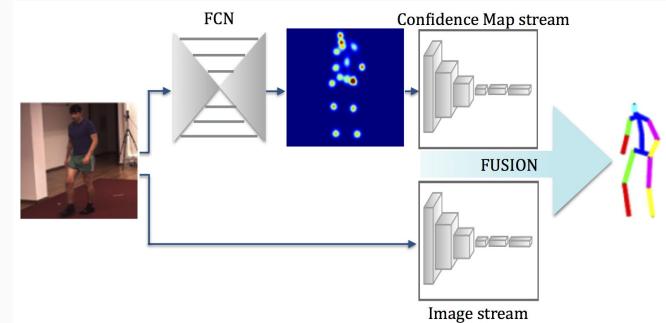
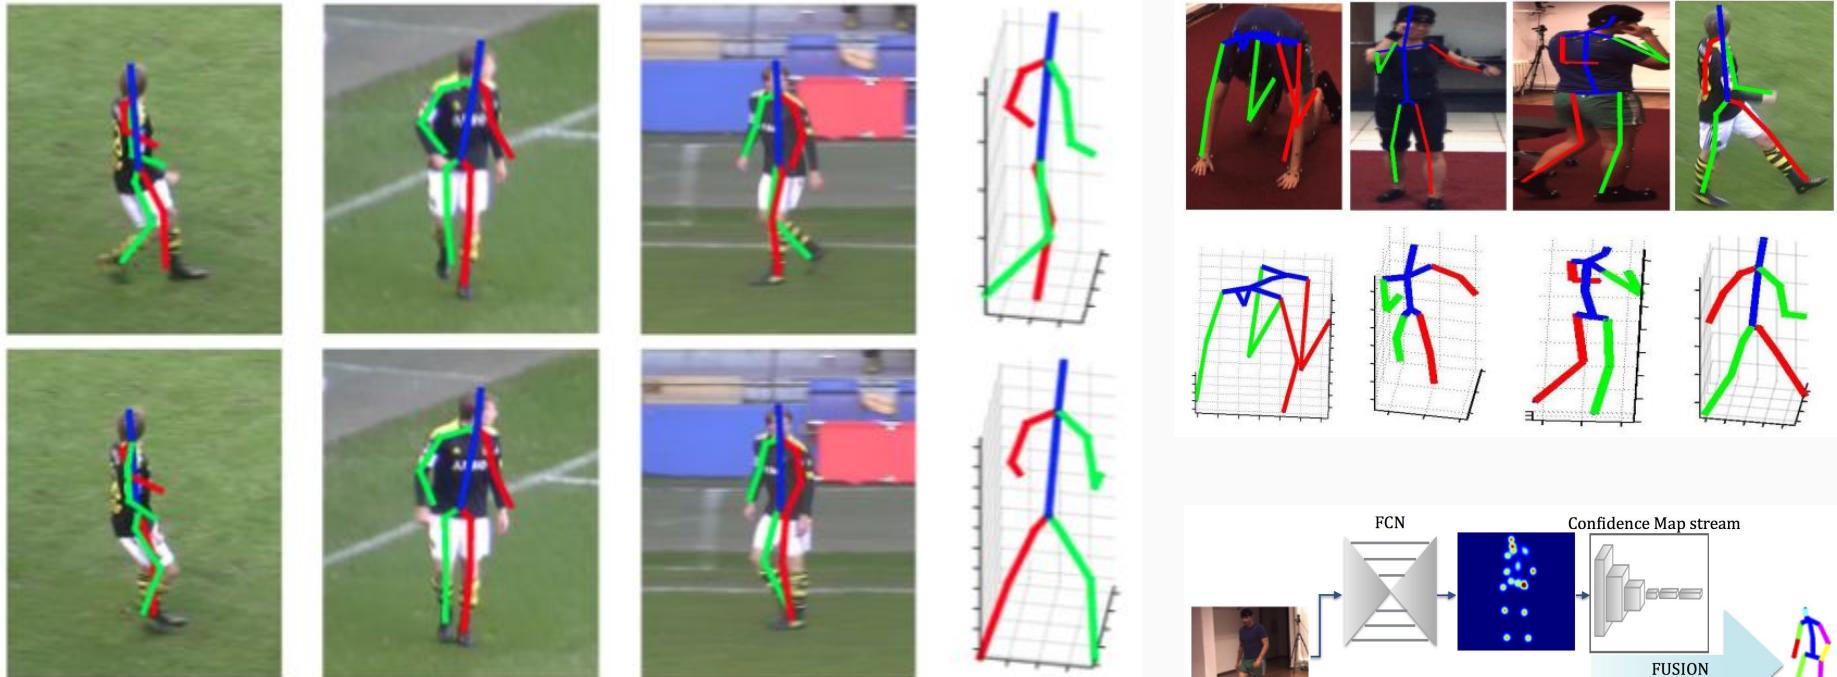
Instance Segmentation

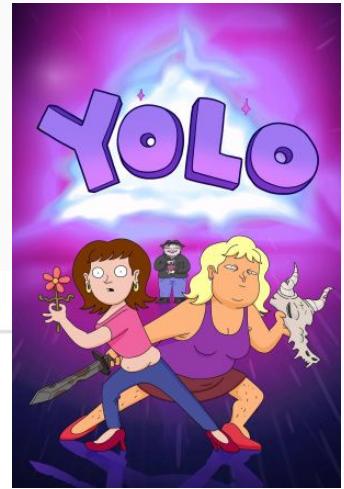


DOG, DOG, CAT

This image is CC0 public domain







## What is YOLO?

- **YOLO (You Only Look Once)**: Real-time object detection architecture.
- **Philosophy**: Single neural network predicts bounding boxes and class probabilities in one pass → extremely fast.
- **Applications**: Detection, segmentation, pose estimation.

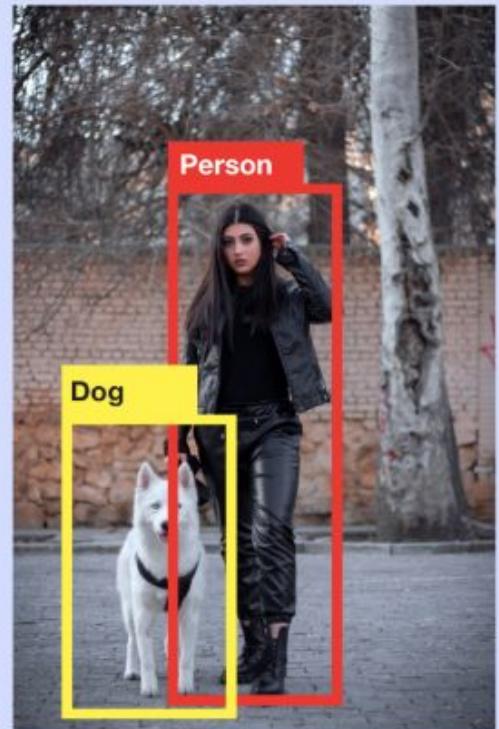
1

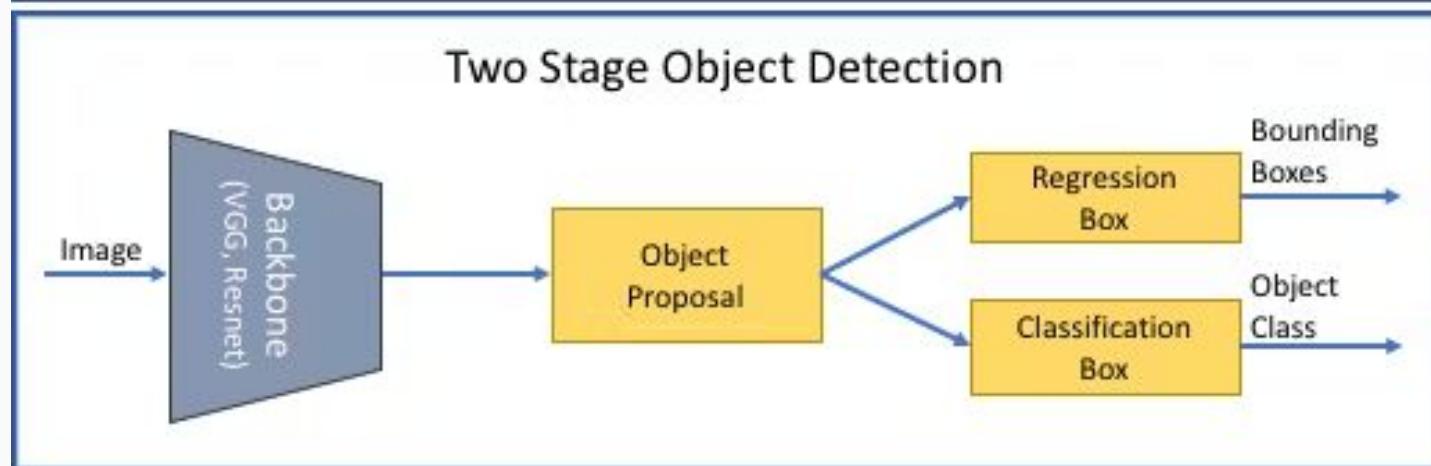
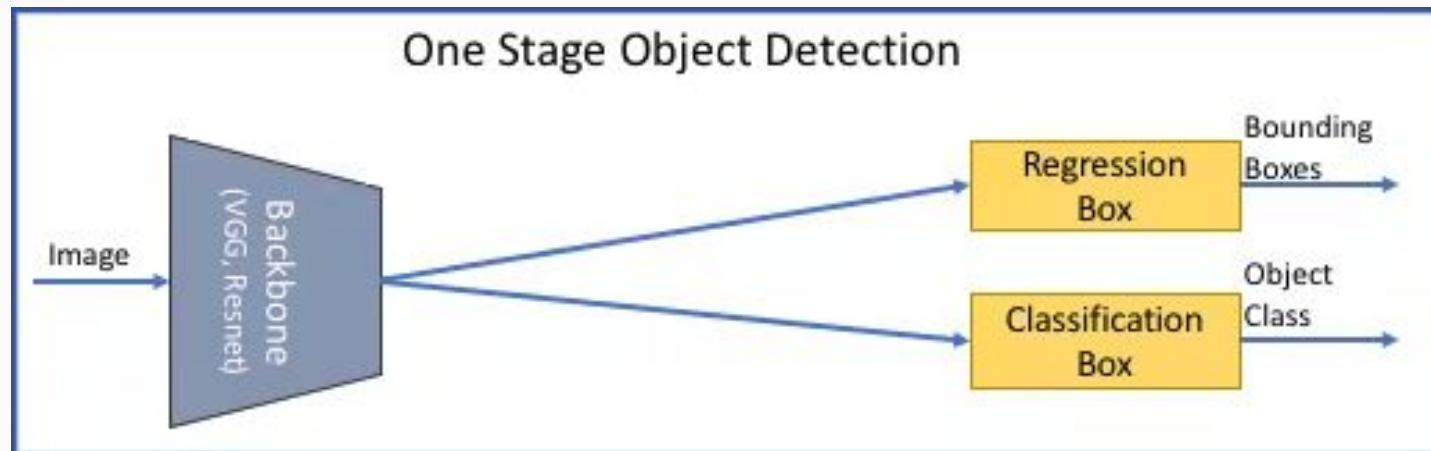


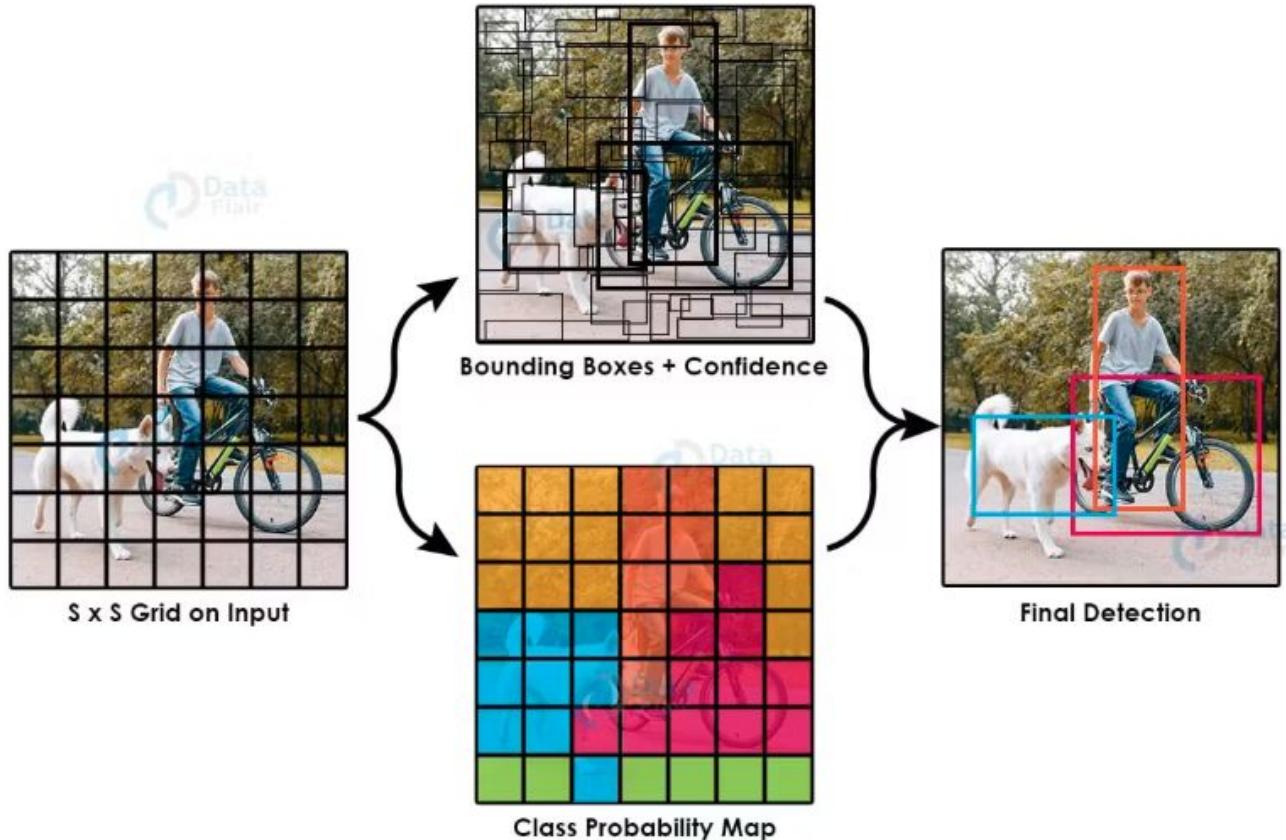
2



3







*Custom Object Detection with YOLO V5*

### Classes

Red Pepper

Aubergine

Broccoli

Lettuce

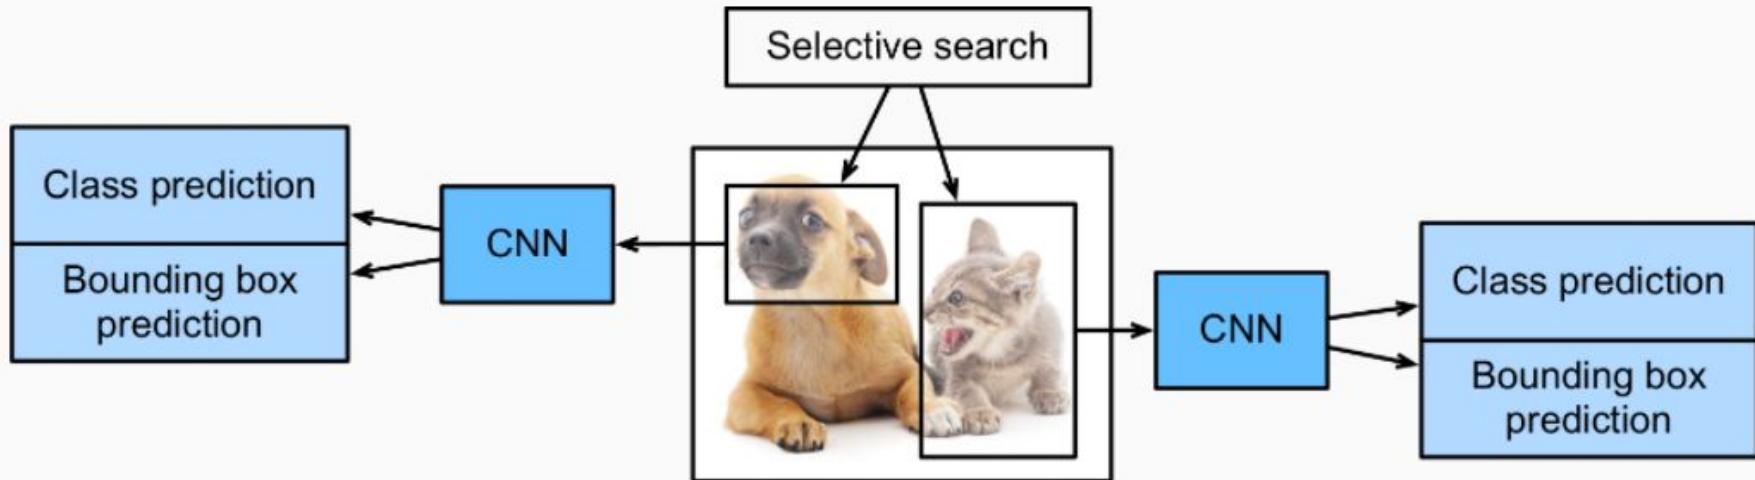
Courgette

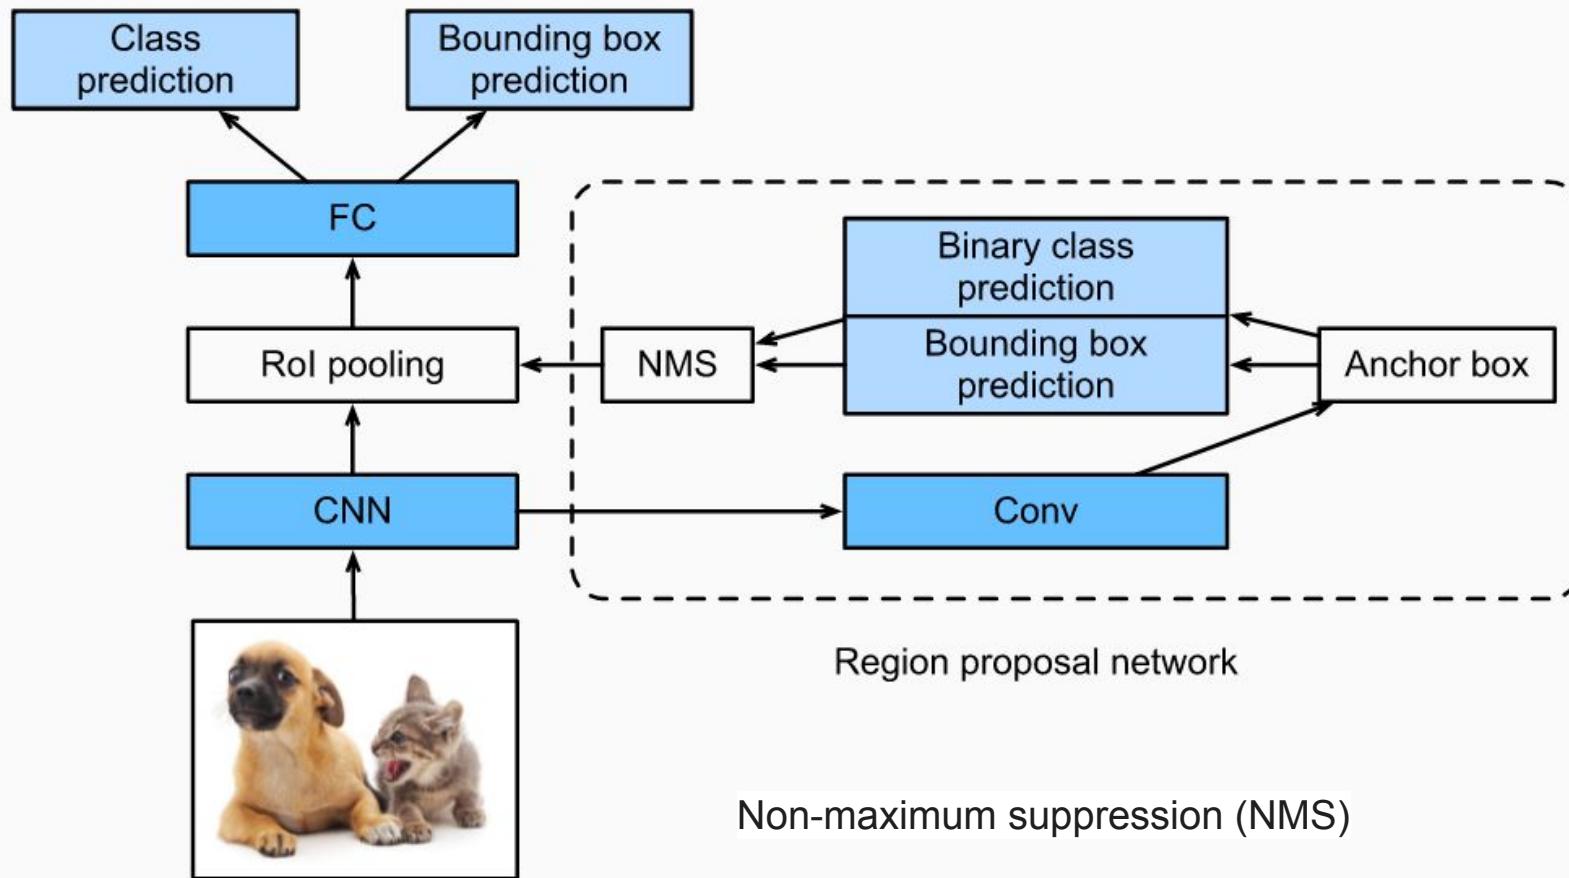
Do products appear fresh?

Instances Frame labels

Classifications (1)







## Non-maximum suppression (NMS)



Applying  
NMS



LearnOpenCV.com

## YOLO Family Overview

- **YOLOv8 Nano (n)**: Smallest, fastest, low computation, suitable for edge devices.
- **YOLOv8 Small (s)**: Lightweight, faster, good for smaller datasets or quick experiments.
- **YOLOv8 Medium (m)**: Balance between speed and accuracy, ideal for most use-cases.
- **YOLOv8 Large (l)**: High accuracy, slower inference, suitable for powerful GPUs.
- **YOLOv8 Extra Large (x)**: Maximum accuracy, large datasets, heavy computation.

**Models**

YOLOv3

YOLOv4

YOLOv5

YOLOv6

YOLOv7

YOLOv8

YOLOv9

YOLOv10

**YOLO11**  NEW

YOLO12

SAM (Segment Anything Model)

SAM 2 (Segment Anything Model 2)

MobileSAM (Mobile Segment Anything Model)

FastSAM (Fast Segment Anything Model)

YOLO-NAS (Neural Architecture Search)

RT-DETR (Realtime Detection Transformer)

YOLO-World (Real-Time Open-Vocabulary Object Detection)

# Ultralytics YOLO11

**Table of contents**[Overview](#)[Key Features](#)[Supported Tasks and Modes](#)[Performance Metrics](#)[Usage Examples](#)[Citations and Acknowledgements](#)[FAQ](#)

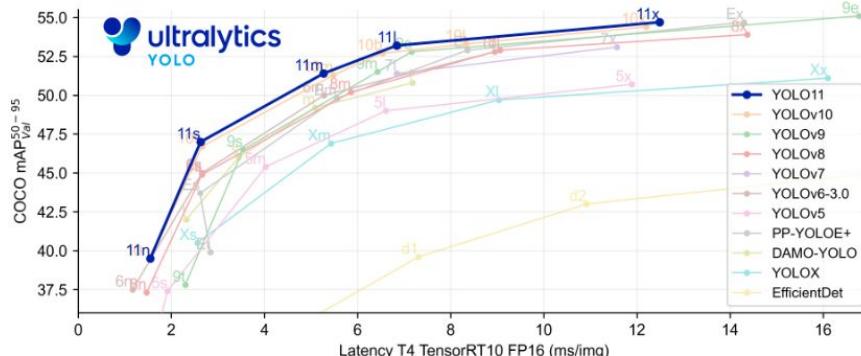
What are the key improvements in Ultralytics YOLO11 compared to previous versions?

How do I train a YOLO11 model for object detection?

What tasks can YOLO11 models perform?

How does YOLO11 achieve greater accuracy with fewer parameters?

Can YOLO11 be deployed on edge devices?

Ask AI 

**Models**

YOLOv3

YOLOv4

YOLOv5

YOLOv6

YOLOv7

YOLOv8

YOLOv9

YOLOv10

**YOLO11  NEW**

YOLO12

SAM (Segment Anything Model)

SAM 2 (Segment Anything Model 2)

MobileSAM (Mobile Segment Anything Model)

FastSAM (Fast Segment Anything Model)

YOLO-NAS (Neural Architecture Search)

RT-DETR (Realtime Detection Transformer)

YOLO-World (Real-Time Open-Vocabulary Object Detection)

YOLOE (Real-Time Seeing Anything)

## Supported Tasks and Modes

YOLO11 builds upon the versatile model range introduced in YOLOv8, offering enhanced support across various computer vision tasks:

Model	Filenames	Task	Inference	Validation	Training	Export
YOLO11	yolo11n.pt yolo11s.pt yolo11m.pt yolo11l.pt yolo11x.pt	Detection	✓	✓	✓	✓
YOLO11-seg	yolo11n-seg.pt yolo11s-seg.pt yolo11m-seg.pt yolo11l-seg.pt yolo11x-seg.pt	Instance Segmentation	✓	✓	✓	✓
YOLO11-pose	yolo11n-pose.pt yolo11s-pose.pt yolo11m-pose.pt yolo11l-pose.pt yolo11x-pose.pt	Pose/Keypoints	✓	✓	✓	✓
YOLO11-obb	yolo11n-obb.pt yolo11s-obb.pt yolo11m-obb.pt yolo11l-obb.pt yolo11x-obb.pt	Oriented Detection	✓	✓	✓	✓
YOLO11-cls	yolo11n-cls.pt yolo11s-cls.pt yolo11m-cls.pt yolo11l-cls.pt yolo11x-cls.pt	Classification	✓	✓	✓	✓

**Table of contents**[Overview](#)[Key Features](#)[Supported Tasks and Modes](#)[Performance Metrics](#)[Usage Examples](#)[Citations and Acknowledgements](#)[FAQ](#)

What are the key improvements in Ultralytics YOLO11 compared to previous versions?

How do I train a YOLO11 model for object detection?

What tasks can YOLO11 models perform?

How does YOLO11 achieve greater accuracy with fewer parameters?

Can YOLO11 be deployed on edge devices?

**Ask AI **

## Performance

**Detection (COCO)**

**Segmentation (COCO)**

**Classification (ImageNet)**

**Pose (COCO)**

**OBB (DOTAv1)**

See [Detection Docs](#) for usage examples with these models trained on **COCO**, which include 80 pre-trained classes.

Model	size (pixels)	mAP <sup>val</sup> 50-95	Speed CPU ONNX (ms)	Speed T4 TensorRT10 (ms)	params (M)	FLOPs (B)
YOLO11n	640	39.5	$56.1 \pm 0.8$	$1.5 \pm 0.0$	2.6	6.5
YOLO11s	640	47.0	$90.0 \pm 1.2$	$2.5 \pm 0.0$	9.4	21.5
YOLO11m	640	51.5	$183.2 \pm 2.0$	$4.7 \pm 0.1$	20.1	68.0
YOLO11l	640	53.4	$238.6 \pm 1.4$	$6.2 \pm 0.1$	25.3	86.9
YOLO11x	640	54.7	$462.8 \pm 6.7$	$11.3 \pm 0.2$	56.9	194.9

## How to Use YOLO

1. Install Ultralytics YOLO (`pip install ultralytics`).
2. Load a pre-trained model or define your own.

python

 Copy  Edit

```
from ultralytics import YOLO  
model = YOLO("yolov8n.pt") # detection
```

3. Train on your dataset with `.train()` method.
4. Evaluate and visualize results with `.val()` and inference methods.



## Example

Python

CLI

PyTorch pretrained `*.pt` models as well as configuration `*.yaml` files can be passed to the `YOLO()` class to create a model instance in Python:

```
from ultralytics import YOLO

# Load a COCO-pretrained YOL011n model
model = YOLO("yolo11n.pt")

# Train the model on the COCO8 example dataset for 100 epochs
results = model.train(data="coco8.yaml", epochs=100, imgsz=640)

# Run inference with the YOL011n model on the 'bus.jpg' image
results = model("path/to/bus.jpg")
```

**Datasets****Detection**

[Argoverse](#)  
[COCO](#)  
[COCO8](#)  
[COCO8-Grayscale](#)  
[COCO8-Multispectral](#)  
[COCO128](#)

[LVIS](#)  
[GlobalWheat2020](#)

[Objects365](#)  
[OpenImagesV7](#)  
[SKU-110K](#)  
[HomeObjects-3K](#)

[VisDrone](#)  
[VOC](#)

[xView](#)

[RF100](#)

[Brain-tumor](#)  
[African-wildlife](#)  
[Signature](#)  
[Medical-pills](#)

**Segmentation**  
[COCO](#)  
[COCO8-seg](#)

## Object Detection

 Back to top

[Bounding box](#) object detection is a computer vision technique that involves detecting and localizing objects in an image by drawing a bounding box around each object.

- [Argoverse](#): A dataset containing 3D tracking and motion forecasting data from urban environments with rich annotations.
- [COCO](#): Common Objects in Context (COCO) is a large-scale object detection, segmentation, and captioning dataset with 80 object categories.
- [LVIS](#): A large-scale object detection, segmentation, and captioning dataset with 1203 object categories.
- [COCO8](#): A smaller subset of the first 4 images from COCO train and COCO val, suitable for quick tests.
- [COCO8-Grayscale](#): A grayscale version of COCO8 created by converting RGB to grayscale, useful for single-channel model evaluation.
- [COCO8-Multispectral](#): A 10-channel multispectral version of COCO8 created by interpolating RGB wavelengths, useful for spectral-aware model evaluation.
- [COCO128](#): A smaller subset of the first 128 images from COCO train and COCO val, suitable for tests.
- [Global Wheat 2020](#): A dataset containing images of wheat heads for the Global Wheat Challenge 2020.
- [Objects365](#): A high-quality, large-scale dataset for object detection with 365 object categories and over 600K annotated images.
- [OpenImagesV7](#): A comprehensive dataset by Google with 1.7M train images and 42k validation images.
- [SKU-110K](#): A dataset featuring dense object detection in retail environments with over 11K images and 1.7 million bounding boxes.
- [HomeObjects-3K](#) : A dataset of annotated indoor scenes featuring 12 common household items, ideal for developing and testing computer vision models in smart home systems, robotics, and augmented reality.

**Table of contents**

[Object Detection](#)  
[Instance Segmentation](#)  
[Pose Estimation](#)  
[Classification](#)  
[Oriented Bounding Boxes \(OBB\)](#)  
[Multi-Object Tracking](#)  
[Contribute New Datasets](#)

[Steps to Contribute a New Dataset](#)  
[Example Code to Optimize and Zip a Dataset](#)

**FAQ**

[What datasets does Ultralytics support for object detection?](#)  
[How do I contribute a new dataset to Ultralytics?](#)  
[Why should I use Ultralytics HUB for my dataset?](#)  
[What are the unique features of Ultralytics YOLO models for computer vision?](#)  
[How can I optimize and zip a dataset using Ultralytics tools?](#)

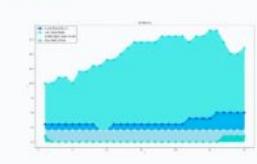
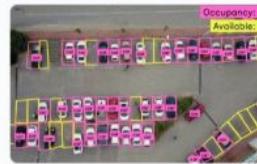


## COCO Format (Common Objects in Context)

- **Definition:** A widely used standard for storing annotations for object detection, segmentation, and keypoint tasks.
- **Structure:** JSON file with fields like:
  - `images` : list of images with `id` , `file_name` , `height` , `width`
  - `annotations` : list of objects with `image_id` , `bbox` (`x`, `y`, `w`, `h`), `category_id` , `segmentation` , `keypoints`
  - `categories` : list of class names with `id`
- **Supports:**
  - Bounding boxes → object detection
  - Segmentation masks → instance segmentation
  - Keypoints → pose estimation

# Ultralytics Solutions: Harness YOLO11 to Solve Real-World Problems

Ultralytics Solutions provide cutting-edge applications of YOLO models, offering real-world solutions like object counting, blurring, and security systems, enhancing efficiency and [accuracy](#) in diverse industries. Discover the power of YOLO11 for practical, impactful implementations.

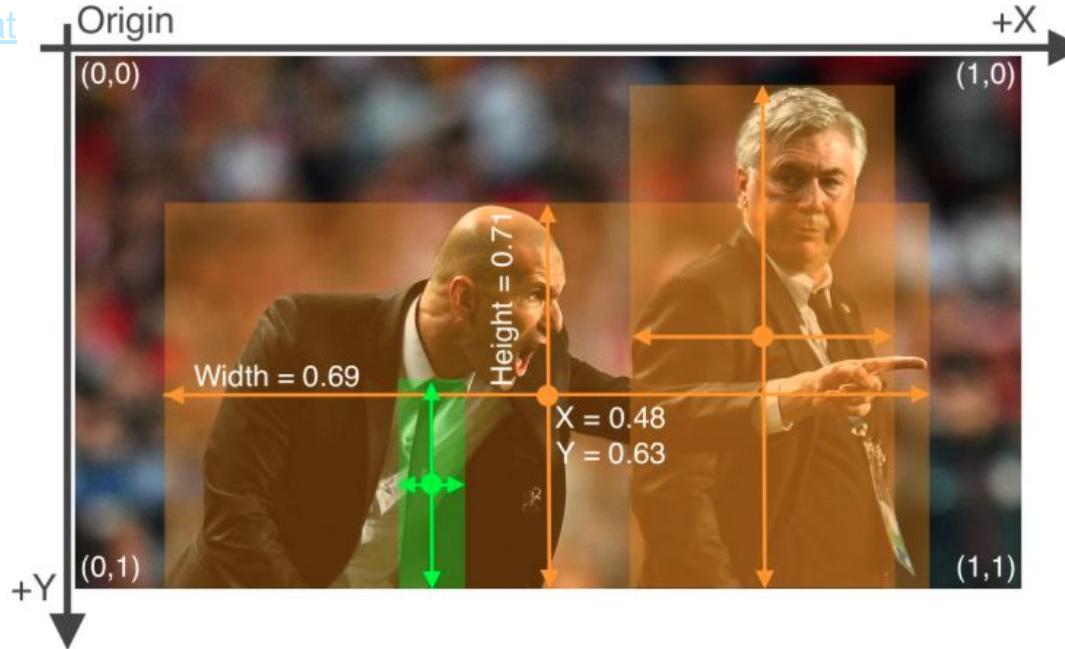


## Ultralytics YOLO format

The Ultralytics YOLO format is a dataset configuration format that allows you to define the dataset root directory, the relative paths to training/validation/testing image directories or `*.txt` files containing image paths, and a dictionary of class names. Here is an example:

```
# Train/val/test sets as 1) dir: path/to/imgs, 2) file: path/to/imgs.txt, or 3) list: [path/to/imgs]
path: coco8 # dataset root dir
train: images/train # train images (relative to 'path') 4 images
val: images/val # val images (relative to 'path') 4 images
test: # test images (optional)

# Classes
names:
  0: person
  1: bicycle
  2: car
  3: motorcycle
  4: airplane
  5: bus
  6: train
  7: truck
  8: boat
  9: traffic light
  10: fire hydrant
  11: stop sign
```



The label file corresponding to the above image contains 2 persons (class 0 ) and a tie (class 27 ):

```
zidane.txt ~
0 0.481719 0.634028 0.690625 0.713278
0 0.741094 0.524306 0.314750 0.933389
27 0.364844 0.795833 0.078125 0.400000
```

coco8			
	Name	Kind	Size
	data.yaml	YAML Document	1 KB
	images	Folder	--
	train	Folder	--
	000000000009.jpg	JPEG image	45 KB
	000000000025.jpg	JPEG image	65 KB
	000000000030.jpg	JPEG image	22 KB
	000000000034.jpg	JPEG image	73 KB
	val	Folder	--
	000000000036.jpg	JPEG image	49 KB
	000000000042.jpg	JPEG image	42 KB
	000000000049.jpg	JPEG image	30 KB
	000000000061.jpg	JPEG image	100 KB
	labels	Folder	--
	train	Folder	--
	000000000009.txt	Plain Text	312 bytes
	000000000025.txt	Plain Text	78 bytes
	000000000030.txt	Plain Text	72 bytes
	000000000034.txt	Plain Text	39 bytes
	val	Folder	--
	000000000036.txt	Plain Text	77 bytes
	000000000042.txt	Plain Text	35 bytes
	000000000049.txt	Plain Text	328 bytes
	000000000061.txt	Plain Text	193 bytes

Macintosh HD > Users > glennjocher > PycharmProjects > datasets > coco8



---

## Preparing Custom Datasets

- **Label Your Own Dataset:**
    - Recommended tool: **Label Studio** → annotate images for detection, segmentation, or keypoints.
  - **Play with Public Datasets:**
    - Recommended tool: **Roboflow** → download, augment, split datasets easily.
-

## Preparing Custom Datasets (Label Your Own)

- Why label your own dataset:
  - Real-world applications often have unique objects not found in public datasets.
  - Proper annotation ensures model learns the right features.
- Recommended Tool: Label Studio
  - Open-source, versatile labeling platform.
  - Supports: bounding boxes, polygons, keypoints, masks, text.
  - Workflow:
    1. Upload images.
    2. Define labeling schema (classes, keypoints).
    3. Annotate images.
    4. Export in YOLO format (`.txt` for detection, masks for segmentation, keypoints for pose).
  - Tips:
    - Start with small batches, check annotation quality.
    - Use keyboard shortcuts to speed up labeling.
    - Consistent labeling improves model performance.
  - Extra Tip for Students:
    - Annotate edge cases (small objects, partially occluded objects) → model generalizes better.
    - Review annotations twice; mislabeled images reduce model accuracy.



# Label Studio

GUIDE HOW TO BUILD AI BENCHMARKS THAT EVOLVE WITH YOUR MODEL



Label Studio

Use Cases

Learn

Docs

Integrations

Enterprise

Search

24,251

Quick Start

# Open Source Data Labeling Platform

The most flexible data labeling platform to fine-tune LLMs, prepare training data, or evaluate AI models.

[Download OSS](#)[Compare Versions](#)

LAST COMMIT: AUGUST 19, 2025 | LATEST VERSION: NIGHTLY

## Quick Start

```
1 # Install the package
# into python virtual environment
2 pip install -U label-studio
3
4 # Launch it!
5 label-studio
```

[PIP](#) [BREW](#) [GIT](#) [DOCKER](#)

## Choose annotation template

Audio classification	Image classification	A	</>	-[]-
Emotion segmentation	Bbox object detection	Text classification	HTML classification	Time Series classification
Speaker diarization	Brush segmentation	Multi classification	HTML NER tagging	Import CSV
Transcription per region	Circular object detector	Named entity recognition	Dialogs & conversations	Import JSON
Transcription whole audio	Keypoints and landmarks	Text summarization	Rate PDF	Segmentation extended
	Polygon segmentation	Word alignment	Rate website	Multi-step annotation
	Multi-image classification		Video classifier	

» Advanced config templates

## Label Studio Playground

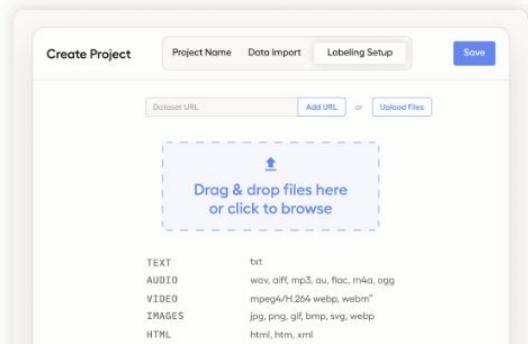
```
1 <View>
2   <Image name="image" value="$image"/>
3   <RectangleLabels name="label" toName="image">
4     <Label value="Airplane" background="green"/>
5     <Label value="Car" background="blue"/>
6   </RectangleLabels>
7 </View>
8
```

Auto Beta

# Label Studio Documentation

Discover our Quick Start installation guide, instructions on building custom UIs, pre-built labeling templates to get you labeling more quickly, and much more.

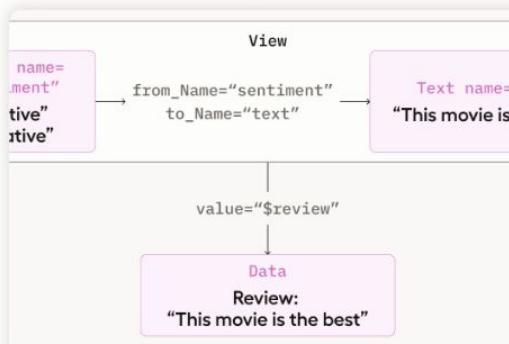
## Quick Start Guide



## Label Studio 101

Brand new to Label Studio? We've created a jam-packed new tutorial with the most important information to get you up and running.

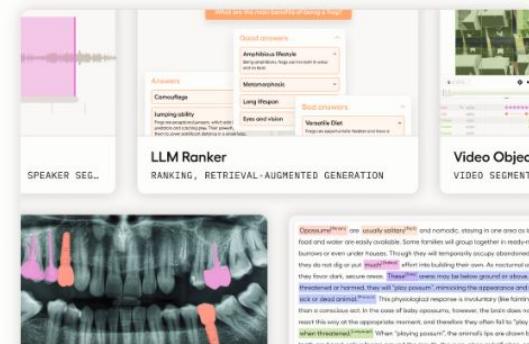
[Get started →](#)



## Machine Learning with Label Studio

Learn the background knowledge and steps required to integrate Machine Learning models into your Label Studio workflow.

[Learn more →](#)



## Use templates

Label Studio provides multiple out-of-the-box labeling configurations to help get you started.

[Explore templates →](#)

# Template Gallery



Runner1<sup>[1]</sup> | Runner2<sup>[2]</sup> | Runner3<sup>[3]</sup> | Runner4<sup>[4]</sup> | Runner5<sup>[5]</sup>

Computer Vision

Object Detection, Semantic Segmentation, Image Classification



Dynamic Labels

Object Detection, Semantic Segmentation, Image Classification



Noun<sup>[1]</sup> | Pronoun<sup>[2]</sup>

I voted for Obama<sup>Noun</sup> because he<sup>Pronoun</sup> was...

Natural Language Processing

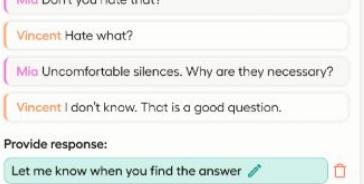
Named Entity Recognition, Text Classification, Relation Extraction



Speaker one<sup>[1]</sup> | Speaker two<sup>[2]</sup>

Audio/Speech Processing

Automatic Speech Recognition, Speaker Segmentation, Intent Classification

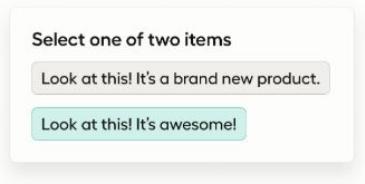


Mia Don't you hate that?  
Vincent Hate what?  
Mia Uncomfortable silences. Why are they necessary?  
Vincent I don't know. That is a good question.

Provide response:  
Let me know when you find the answer

Conversational AI

Natural Language Understanding, Chatbot Response Generation, Slot Filling



Select one of two items

Look at this! It's a brand new product.  
Look at this! It's awesome!

Ranking & Scoring

Pairwise Classification, Document Retrieval

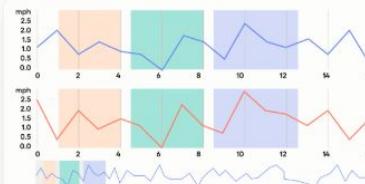


This is a movie that delivers everything almost right in your face.

Alice	<input type="radio"/>	5.0	<input type="radio"/>
Bob	<input type="radio"/>	4.8	<input type="radio"/>
Carl	<input type="radio"/>	3.7	<input type="radio"/>

Structured Data Parsing

Freeform Metadata, Tabular Data, Pdf Classification



Run<sup>[1]</sup> | Walk<sup>[2]</sup> | Fly<sup>[3]</sup> | Swim<sup>[4]</sup> | Ride<sup>[5]</sup>

Time Series Analysis

Activity Recognition, Forecasting, Outliers And Anomaly Detection

# Inventory Tracking

Inventory Tracking system allows you to label exact products by given brand names illustrated by relevant sample photo. Every task with shelf photo also has a list of assiated brands to label.



## Using Public Datasets (Roboflow)

- Why use public datasets:
  - Save time & effort for training and experimentation.
  - Standard datasets allow benchmarking and comparison.
- Recommended Tool: Roboflow
  - Upload or import public datasets (COCO, Pascal VOC, Kaggle, GitHub).
  - Provides:
    - Auto-splitting into train/val/test sets.
    - Data augmentation: flips, rotation, brightness, noise.
    - Conversion into YOLO format automatically\*\*.
  - Steps:
    1. Import dataset or upload your images.
    2. Apply augmentation if needed.
    3. Export dataset in YOLOv8-ready format.
  - Tips:
    - Explore multiple public datasets → combine to increase variety.
    - Use preview function to verify images & annotations before training.
- Extra Tip for Students:
  - Use Roboflow for quick prototyping before investing time in labeling your own data.
  - Augment dataset to handle class imbalance or improve robustness.

# Everything you need to build and deploy computer vision applications.

Used by over 1 million engineers to create datasets, train models, and deploy to production.

Get Started

Request a Demo

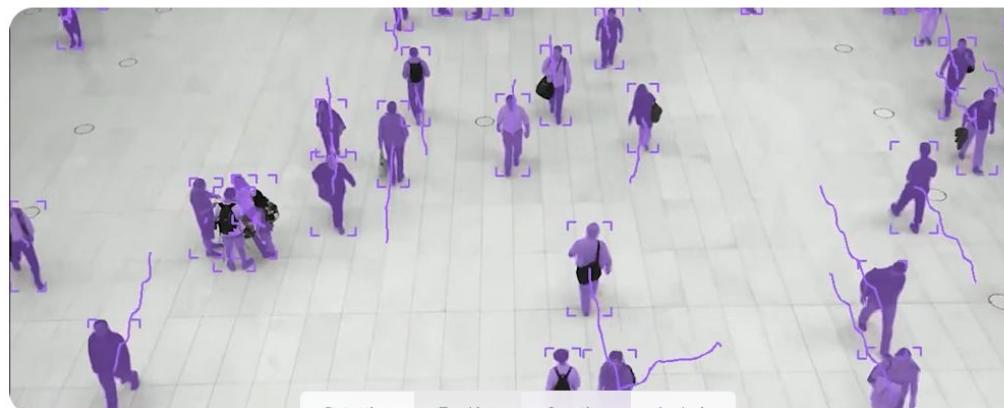


Over 16,000 organizations build with Roboflow. [See our customers >](#)

[Products](#) [Solutions](#) [Resources](#) [Pricing](#) [Docs](#) [Blog](#)[Sign In](#)[Book a demo](#)[Get Started](#)

# Everything you need to build and deploy computer vision applications.

Used by over 1 million engineers to create datasets, train models, and deploy to production.

[Get Started](#)[Request a Demo](#)

## PLATFORM

**Universe**

Open source computer vision datasets and pre-trained models

**Annotate**

Label images fast with AI-assisted data annotation

**Train**

Hosted model training infrastructure and GPU access

**Workflows**

Low-code interface to build pipelines and applications

**Deploy**

Run models on device, at the edge, in your VPC, or via API

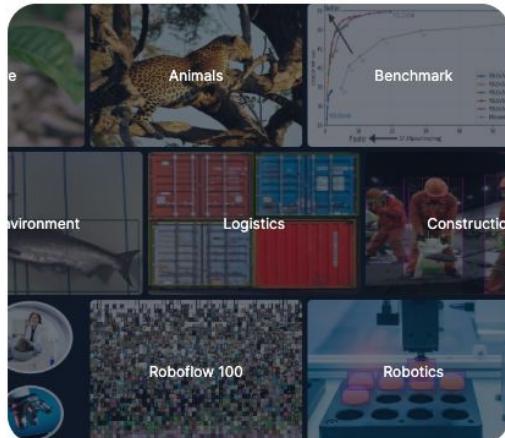
# You need to better vision a

Engineers to create datasets, t

**Get Started****Request**

# Search 575,000+ Datasets for Any Use Case

Explore and analyze datasets containing 750 million+ images



Collections of datasets and pre-trained models for specific use cases

Search and filter for relevant projects to find the best resources



Evaluate the health of datasets and accuracy of models before using them

## PROJECTS (100)



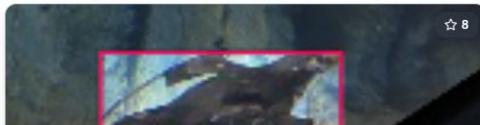
Object Detection Model yolov5

**solar panels**  
by Roboflow 100161 images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**soccer players**  
by Roboflow 100163 images • 1 model  
Updated 2 years ago

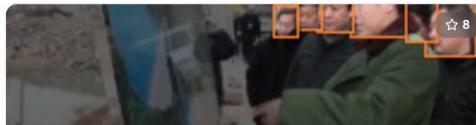
Object Detection Model yolov5

**halo infinite angel videogame**  
by Roboflow 100669 images • 1 model  
Updated a year ago

Object Detection Model yolov5

**aquarium**  
by Roboflow 100638 images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**furniture**  
by Roboflow 100689 images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**street work**  
by Roboflow 100873 images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**leaf disease**  
by Roboflow 1002501 Images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**cotton**  
by Roboflow 100406 images • 1 model  
Updated 2 years ago

Object Detection Model yolov5

**chess pieces**  
by Roboflow 100289 Images • 1 model  
Updated 2 years ago

Search...

## Filter By Task

- [All Models](#)
- Object Detection
- Classification
- Instance Segmentation
- Semantic Segmentation
- Keypoint Detection
- Vision Language
- OCR
- Pose Estimation
- Chart Question Answering
- Document Question Answering (DocQA)
- Video Classification
- Open Vocabulary Object Detection
- Multi-Label Classification
- Region Proposal
- Phrase Grounding
- Referring Expression Segmentation
- Zero Shot Segmentation

## Filter By Feature

- [Foundation Vision](#)
- [Multimodal Vision](#)
- [LLMs with Vision Capabilities](#)



## Model Playground

Compare VLM Models Side-by-Side

## RF-DETR

RF-DETR is a SOTA, real-time object detection model architecture developed by Roboflow and released under the Apache 2.0 license.

## OBJECT DETECTION

Apache 2.0 license • Released Mar 20, 2025

[View Model Details](#)[Deploy with free GPU](#)

## YOLOv12

YOLOv12 is a state-of-the-art computer vision model you can use for detection, segmentation, and more.

## OBJECT DETECTION

700 stars • Released Feb 18, 2025

[View Model Details](#)

## Segment Anything 2

Segment Anything 2 (SAM 2) is a real-time image and video segmentation model.

## INSTANCE SEGMENTATION

3300 stars • Apache 2.0 license • Released Jul 29, 2024

[View Model Details](#)

## GPT-4o

GPT-4o is OpenAI's third major iteration of GPT-4 expanding on the capabilities of GPT-4 with Vision

## VISION LANGUAGE

Released May 13, 2024

## PaliGemma

PaliGemma is a vision language model (VLM) by Google that has multimodal capabilities.

VISION LANGUAGE [DEPLOY WITH ROBOFLOW](#) 2.0k+ stars • Custom Google license  
Released May 14, 2024

## YOLOv12

YOLOv12 is a state-of-the-art computer vision model you can use for detection, segmentation, and more.

OBJECT DETECTION

700 stars • Released Feb 18, 2025

[View Model Details](#)

## GPT-4o

GPT-4o is OpenAI's third major iteration of GPT-4 expanding on the capabilities of GPT-4 with Vision

VISION LANGUAGE

Released May 13, 2024

[View Model Details](#)

## YOLOv9

YOLOv9 is an object detection model architecture released on February 21st, 2024.

OBJECT DETECTION

DEPLOY WITH ROBOFLOW

0 stars • GPL-3.0 license • Released Feb 21, 2024

## Segment Anything 2

Segment Anything 2 (SAM 2) is a real-time image and video segmentation model.

INSTANCE SEGMENTATION

3300 stars • Apache 2.0 license • Released Jul 29, 2024

[View Model Details](#)

## PaliGemma

PaliGemma is a vision language model (VLM) by Google that has multimodal capabilities.

VISION LANGUAGE

DEPLOY WITH ROBOFLOW

0 stars • 2.0k+ stars • Custom Google license

Released May 14, 2024

[View Model Details](#)

## YOLO-World

YOLO-World is a zero-shot object detection model.

OBJECT DETECTION

DEPLOY WITH ROBOFLOW

0 stars • 2.9k+ stars • GPL-3.0 license • Released Jan 31, 2024

## YOLOv8

YOLOv8 is a state-of-the-art object detection and image segmentation model created by Ultralytics, the developers of YOLOv5.

OBJECT DETECTION DEPLOY WITH ROBOFLOW

0 stars • 21.1k+ stars • AGPL-3.0 license • Released Jan 10, 2023

[View Model Details](#)

Deploy with free GPU

## GPT-5

GPT-5 is a multimodal language model developed by OpenAI.

VISION LANGUAGE

Released Aug 7, 2025

[View Model Details](#)

## Depth Anything V2

Depth-Anything-V2 is a depth estimation model developed by researchers from HKU and TikTok.

0 stars • 6100 stars • Apache 2.0 license

[View Model Details](#)

## GPT-4.1

GPT-4.1 is a multimodal model developed by OpenAI that comes in three sizes: GPT-4.1, mini, and nano.

Released Apr 14, 2025

[View Model Details](#)

## Claude 3.7 Sonnet

Claude 3.7 is a multimodal "hybrid reasoning" model developed by Anthropic.

Released Feb 24, 2024

## Phi-4 Multimodal

Phi-4 Multimodal is a multimodal language model developed by Microsoft.

VISION LANGUAGE

0 stars • MIT License

## SmolVLM2

SmolVLM2 is a multimodal image and video understanding model developed by engineers on the Hugging Face TB (Textbook) Research team.

0 stars • 2000 stars • Apache 2.0 license • Released Feb 20, 2025

[View Model Details](#)

## Moondream 2

Moondream 2 is the latest model in the Moondream series of "tiny vision language models".

0 stars • 7600 stars • Apache 2.0 license

[View Model Details](#)

## Gemma 3

Gemma 3 is a multimodal language model developed by Google.

Gemma License license • Released Mar 11, 2025

[View Model Details](#)

## OpenAI o3-mini

OpenAI o3-mini is a multimodal reasoning model developed by OpenAI.

VISION LANGUAGE

Released Jan 31, 2025

[View Model Details](#)

## Qwen2.5-VL

Qwen2.5-VL is a multimodal vision-language model developed by the Qwen team at Alibaba Cloud.

VISION LANGUAGE

8600 stars • Apache 2.0 license • Released Jan 28, 2025

[View Model Details](#)

## PalGemma-2

PalGemma-2 is a multimodal model developed by Google.

VISION LANGUAGE DEPLOY WITH ROBOFLOW

Gemma License license • Released Dec 5, 2024

[View Model Details](#)

## Combining Both Approaches

- **Workflow:**
  1. Start with public dataset (Roboflow) to experiment with YOLO families (n, s, m, l, x).
  2. Label your own custom dataset (Label Studio) for your project-specific objects.
  3. Merge datasets if necessary → improve generalization.
  4. Train YOLO models using both datasets → compare performance.
- **Student Advice:**
  - Keep dataset organized: /images/train , /images/val , /labels/train , /labels/val .
  - Ensure .yaml config paths match dataset folder structure.
  - Always validate dataset visually before training → catch annotation mistakes early.

## Tricks to Revise and Improve YOLO

- **Training Options:** epochs, batch size, image size, optimizer.
- **Augmentation:** Mosaic, flips, color jitter, rotation.
- **Fine-Tuning:** Start from pre-trained weights → adapt to your dataset.
- **Handling Imbalanced Classes:** Focal loss, class weights, oversampling.
- **Custom Loss Functions:** If default loss isn't sufficient for edge cases.
- **Experimentation Tricks:**
  - Mix different YOLO families (n, s, m, l, x).
  - Advanced augmentation strategies.
  - Add dropout, label smoothing, or additional heads for segmentation/pose.

## Lab Today (Hands-On)

Students will run 5 in-class notebooks:

1. **Brain Tumor Detection** – binary detection
2. **Medical Pills Detection** – binary detection
3. **African Wildlife Detection** – multi-class detection
4. **Crack Segmentation** – semantic segmentation
5. **Tiger Pose Estimation** – keypoint detection

Name ↑

- Week10\_InClass\_01\_BrianTumorDetection.ipynb ★
- Week10\_InClass\_02\_MedicalPillsDetection.ipynb ★
- Week10\_InClass\_03\_AfricanWildlifeDetection.ipynb ★
- Week10\_InClass\_04\_CrackSegmentation.ipynb ★
- Week10\_InClass\_05\_TigerPose.ipynb ★

Name ↑

 Week10\_InClass\_01\_BrianTumorDetection.ipynb ★



## Homework

- **Dataset:** Use **Brain Tumor Detection dataset** (Week10\_InClass\_01).
- **Task:** Run YOLOv8 models from all families: n, s, m, l, x.
- **Goal:** Find which model performs best.
- **Bonus:** Apply tricks like focal loss, advanced augmentation, or fine-tuning to beat larger models → earn extra points.